

MedCLIP Zero-Shot Disease Classification
IU-CXR sanity subset (pathology-focused)

Tag / Image variant: raw
Total images in CSV: 29

Ground truth source:
- chexpert_label (single-label per image)

Evaluation setup:

- For each image, MedCLIP PromptClassifier (CheXpert head) produces probabilities for 5 diseases:
Atelectasis, Cardiomegaly, Consolidation, Edema, Pleural Effusion.
- This script converts those scores into disease-present vs absent labels using a 0.5 decision threshold and computes per-disease metrics.

Micro-averaged performance across all 5 diseases (threshold 0.5):

- Micro Accuracy : 0.138
- Micro Precision: 0.138
- Micro Recall : 1.000
- Micro F1 : 0.242
- Total label decisions evaluated: 145

Top-K disease ranking (only images with ≥ 1 positive label):

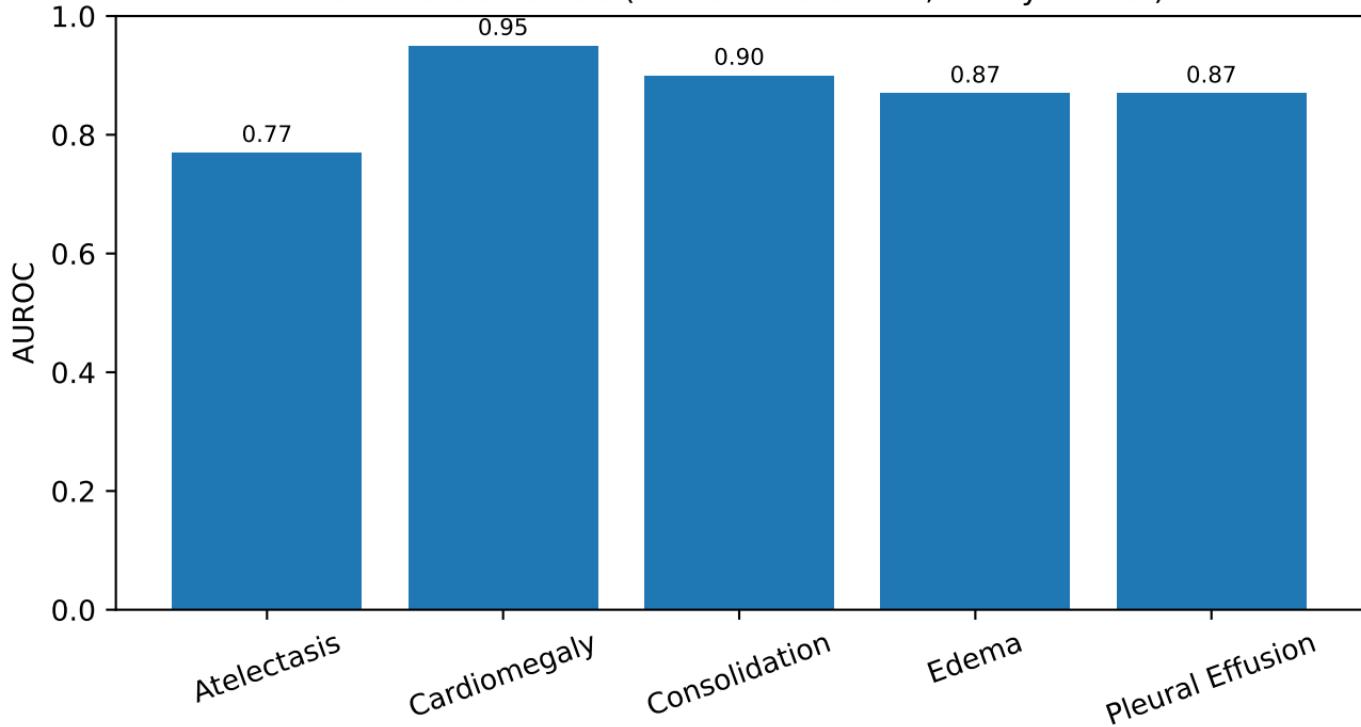
- Top-1 accuracy: 0.600 (N=20)
 - Top-3 accuracy: 0.950 (N=20)
- (Top-K is counted correct if any true disease appears in the top-K scores.)

Any-hit metric (at least one disease present & predicted):

We reduce each image to a single binary decision:

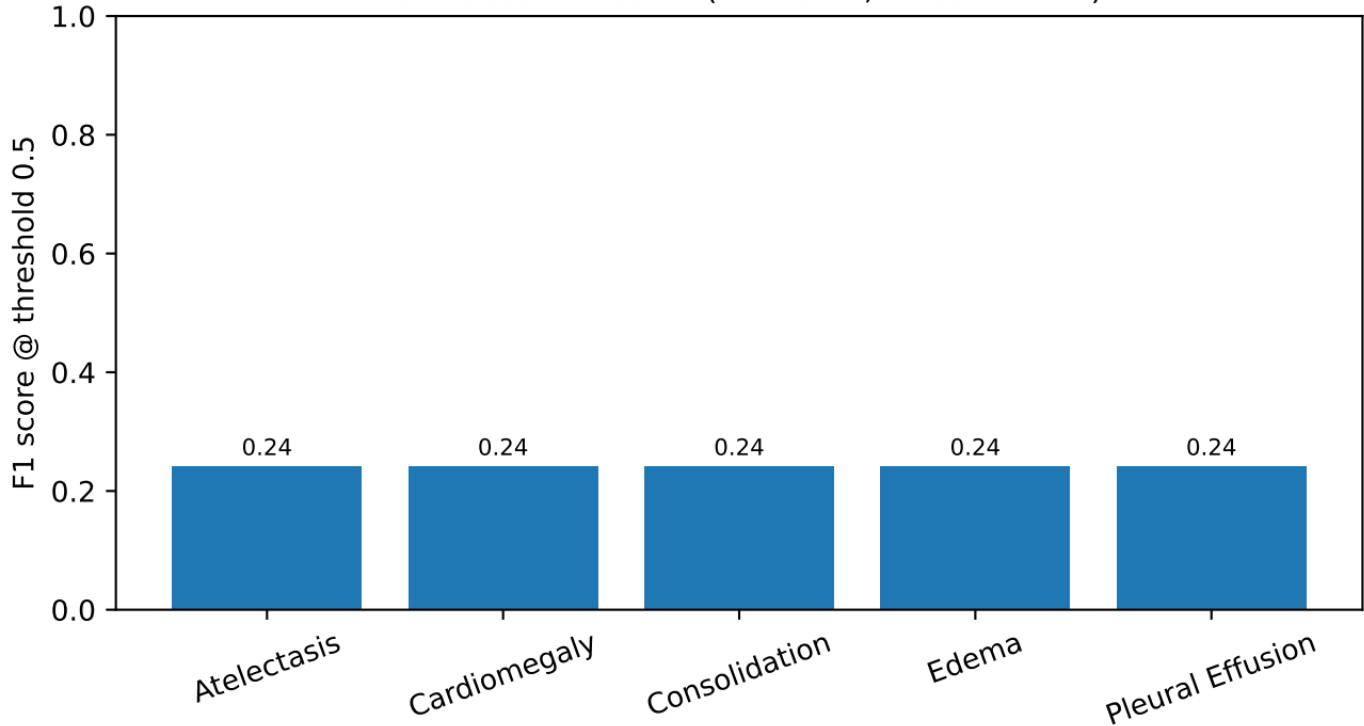
- GT ≥ 1 of the 5 diseases present vs none,
Pred ≥ 1 of the 5 diseases positive vs none.
- Any-hit Accuracy : 0.690
 - Any-hit Precision: 0.690
 - Any-hit Recall : 1.000
 - Any-hit F1 : 0.816

Per-disease AUROC (MedCLIP zero-shot, sanity subset)



AUROC (Area Under ROC Curve) measures how well the model ranks positive cases above negatives for each disease. Values closer to 1.0 indicate better ranking performance, regardless of any fixed threshold.

Per-disease F1 score (zero-shot, threshold 0.5)



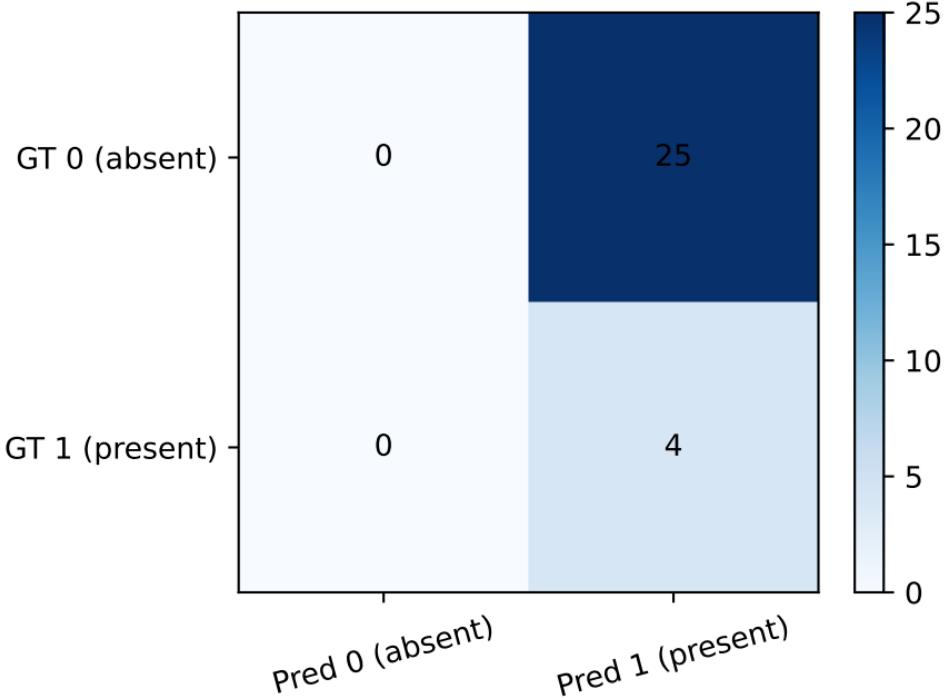
F1 score combines precision and recall for each disease at the 0.5 probability threshold.
It is especially useful when classes are imbalanced (few positives vs many negatives).

Per-disease zero-shot metrics (all images in this subset)

| Disease | Positives | Negatives | AUROC | AUPRC | Acc@0.5 | Prec@0.5 | Rec@0.5 | F1@0.5 | TP | FP | TN | FN |
|------------------|-----------|-----------|-------|-------|---------|----------|---------|--------|----|----|----|----|
| Atelectasis | 4 | 25 | 0.770 | 0.353 | 0.138 | 0.138 | 1.000 | 0.242 | 4 | 25 | 0 | 0 |
| Cardiomegaly | 4 | 25 | 0.950 | 0.812 | 0.138 | 0.138 | 1.000 | 0.242 | 4 | 25 | 0 | 0 |
| Consolidation | 4 | 25 | 0.900 | 0.600 | 0.138 | 0.138 | 1.000 | 0.242 | 4 | 25 | 0 | 0 |
| Edema | 4 | 25 | 0.870 | 0.412 | 0.138 | 0.138 | 1.000 | 0.242 | 4 | 25 | 0 | 0 |
| Pleural Effusion | 4 | 25 | 0.870 | 0.415 | 0.138 | 0.138 | 1.000 | 0.242 | 4 | 25 | 0 | 0 |

Each row summarizes MedCLIP's zero-shot behavior for one disease. We report ranking metrics (AUROC, AUPRC) and threshold-based metrics (Accuracy, Precision, Recall, F1 at 0.5).

Atelectasis: confusion matrix (threshold = 0.5)

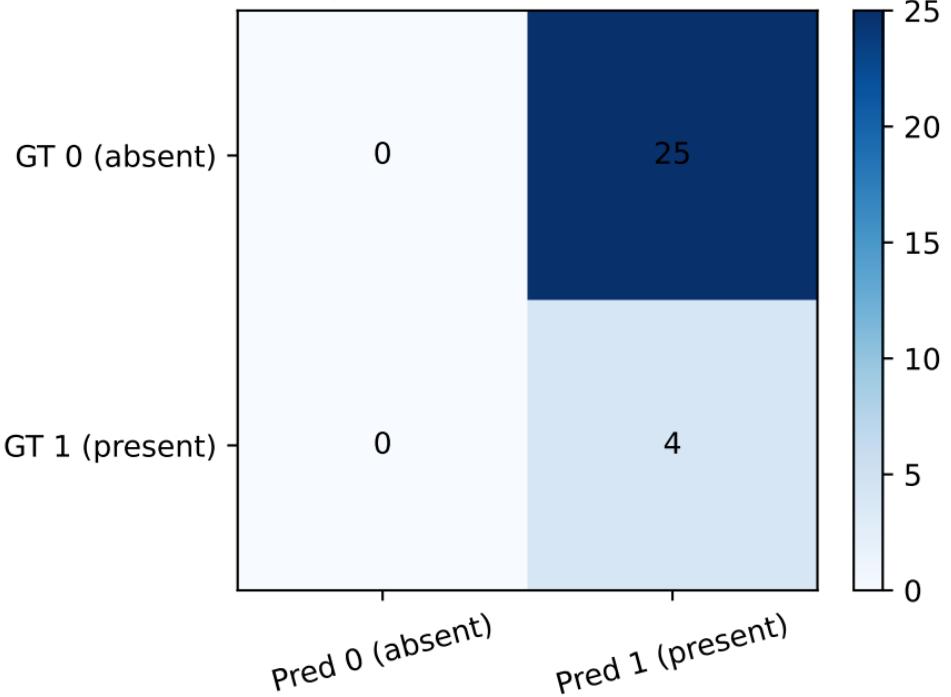


For Atelectasis, positives = 4, negatives = 25.

Interpretation:

- TN (top-left): correctly predicted as no disease.
- FP (top-right): predicted disease but GT is negative.
- FN (bottom-left): missed true disease cases.
- TP (bottom-right): correctly detected disease.

Cardiomegaly: confusion matrix (threshold = 0.5)

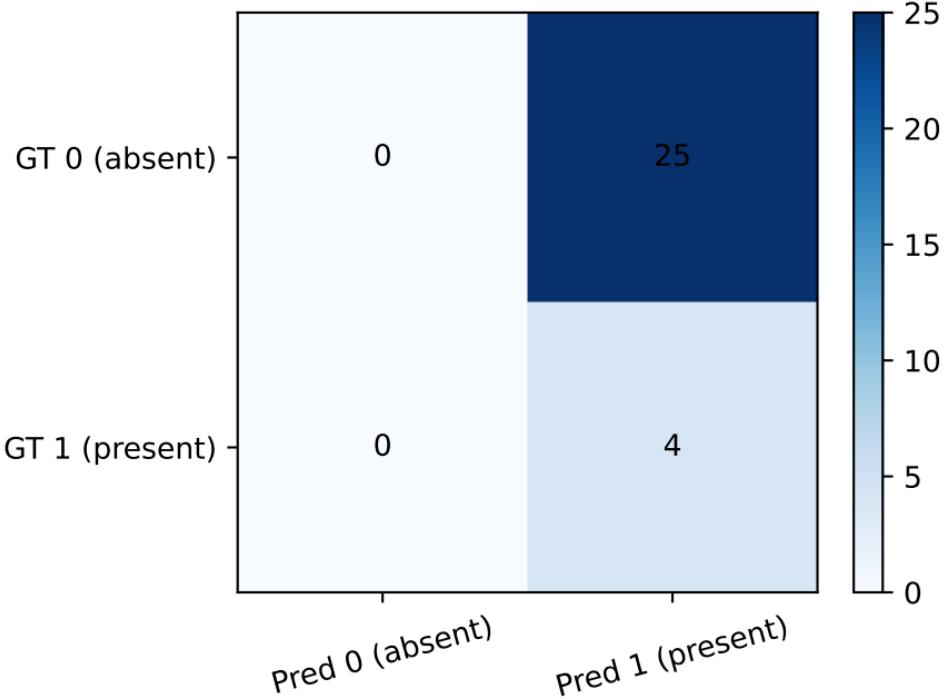


For Cardiomegaly, positives = 4, negatives = 25.

Interpretation:

- TN (top-left): correctly predicted as no disease.
- FP (top-right): predicted disease but GT is negative.
- FN (bottom-left): missed true disease cases.
- TP (bottom-right): correctly detected disease.

Consolidation: confusion matrix (threshold = 0.5)

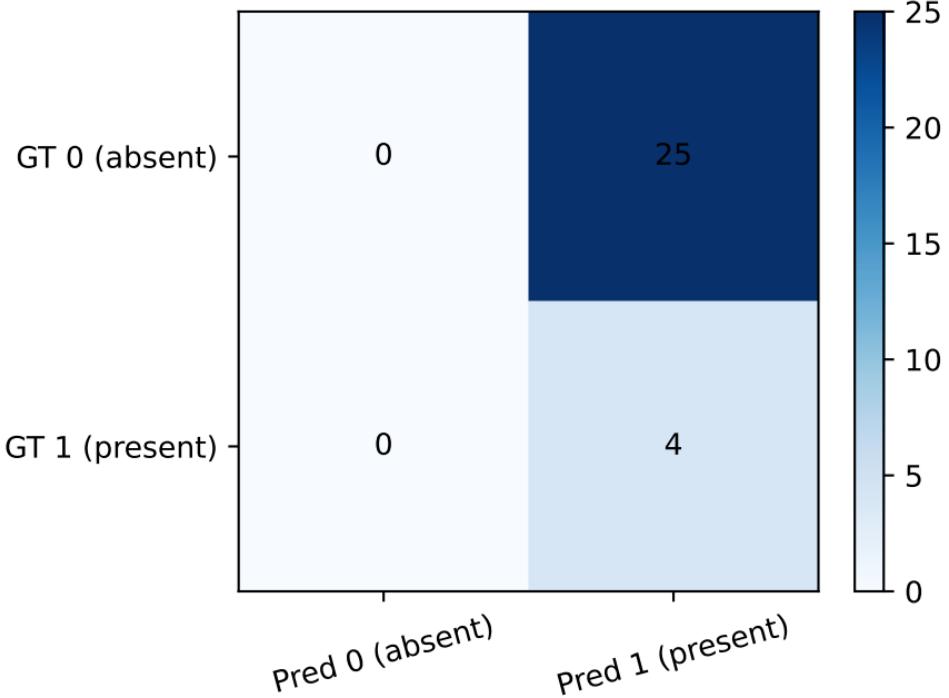


For Consolidation, positives = 4, negatives = 25.

Interpretation:

- TN (top-left): correctly predicted as no disease.
- FP (top-right): predicted disease but GT is negative.
- FN (bottom-left): missed true disease cases.
- TP (bottom-right): correctly detected disease.

Edema: confusion matrix (threshold = 0.5)

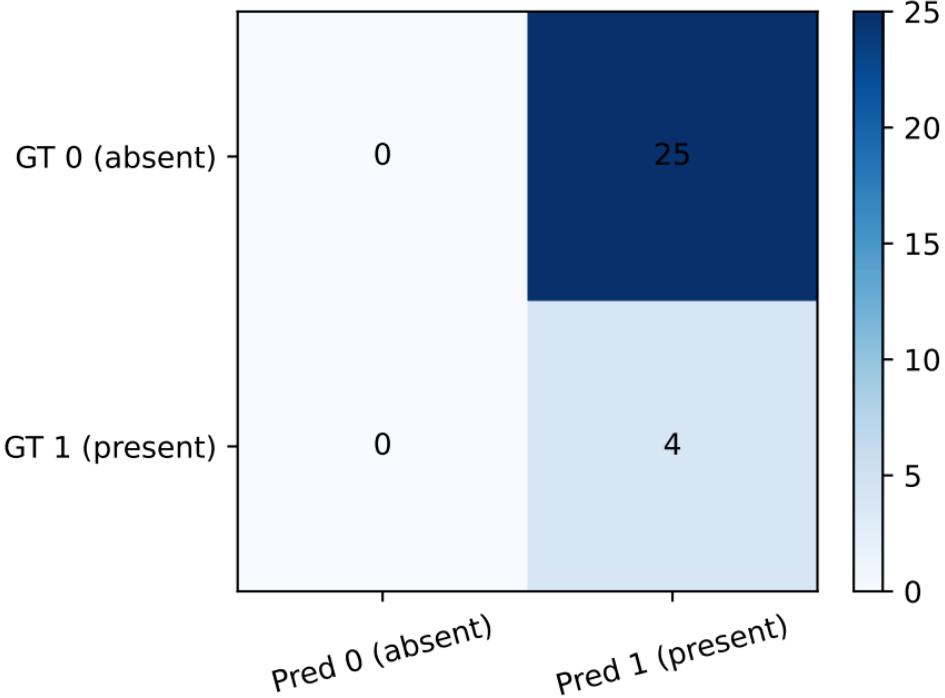


For Edema, positives = 4, negatives = 25.

Interpretation:

- TN (top-left): correctly predicted as no disease.
- FP (top-right): predicted disease but GT is negative.
- FN (bottom-left): missed true disease cases.
- TP (bottom-right): correctly detected disease.

Pleural Effusion: confusion matrix (threshold = 0.5)

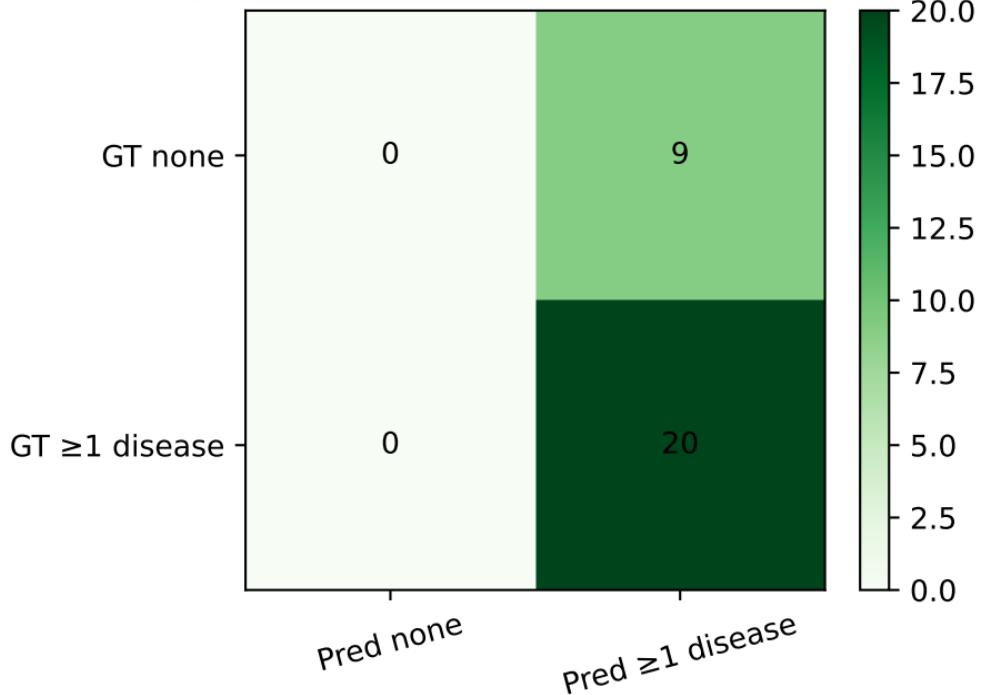


For Pleural Effusion, positives = 4, negatives = 25.

Interpretation:

- TN (top-left): correctly predicted as no disease.
- FP (top-right): predicted disease but GT is negative.
- FN (bottom-left): missed true disease cases.
- TP (bottom-right): correctly detected disease.

Any-disease confusion matrix (threshold = 0.5)



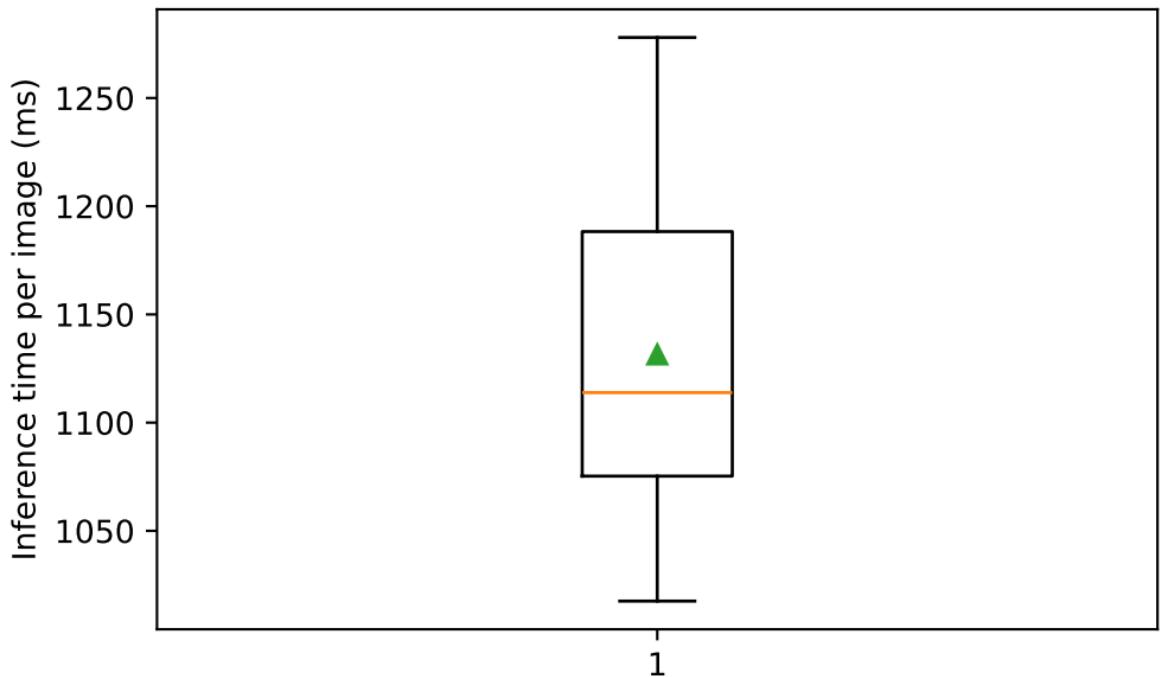
We compress the multi-label task into a single 'any disease' decision:

- GT ≥ 1 disease vs GT none
- Pred ≥ 1 disease vs Pred none (prob ≥ 0.5).

TN=0, FP=9, FN=0, TP=20

Accuracy=0.690, Precision=0.690, Recall=1.000, F1=0.816.

MedCLIP zero-shot disease head - per-image latency



Boxplot summarizing median, interquartile range, mean and outliers for per-image inference time. This is useful to compare runtime across different image variants.

Per-image labels & scores (rows 1-15 of 29)

| filename | chexpert | Atelectasi | Atelectasi | Cardiome | Cardiome | Consolida | Consolida | Edema_gt | Edema_p | Pleural Eft | Pleural Effusion_p |
|-----------------|-------------|------------|------------|----------|----------|-----------|-----------|----------|---------|-------------|--------------------|
| 3638_IM-1804-10 | Cardiome | 0 | 0.547 | 1 | 0.583 | 0 | 0.549 | 0 | 0.572 | 0 | 0.538 |
| 2716_IM-1181-20 | Consolida | 0 | 0.609 | 0 | 0.622 | 1 | 0.591 | 0 | 0.688 | 0 | 0.745 |
| 2422_IM-0965-20 | Consolida | 0 | 0.577 | 0 | 0.570 | 1 | 0.592 | 0 | 0.568 | 0 | 0.593 |
| 3249_IM-1539-20 | Edema | 0 | 0.607 | 0 | 0.603 | 0 | 0.574 | 1 | 0.616 | 0 | 0.780 |
| 1439_IM-0282-10 | Atelectasi | 1 | 0.563 | 0 | 0.563 | 0 | 0.544 | 0 | 0.546 | 0 | 0.575 |
| 1439_IM-0282-10 | Atelectasi | 1 | 0.584 | 0 | 0.562 | 0 | 0.576 | 0 | 0.539 | 0 | 0.616 |
| 3249_IM-1539-10 | Edema | 0 | 0.615 | 0 | 0.637 | 0 | 0.562 | 1 | 0.627 | 0 | 0.762 |
| 91_IM-2415-2001 | Pneumoth | 0 | 0.555 | 0 | 0.548 | 0 | 0.541 | 0 | 0.542 | 0 | 0.545 |
| 3396_IM-1640-30 | Pleural Eft | 0 | 0.612 | 0 | 0.544 | 0 | 0.623 | 0 | 0.543 | 1 | 0.760 |
| 2502_IM-1027-10 | Pleural Eft | 0 | 0.561 | 0 | 0.565 | 0 | 0.615 | 0 | 0.601 | 1 | 0.708 |
| 2716_IM-1181-10 | Consolida | 0 | 0.572 | 0 | 0.564 | 1 | 0.583 | 0 | 0.681 | 0 | 0.783 |
| 2422_IM-0965-10 | Consolida | 0 | 0.574 | 0 | 0.552 | 1 | 0.627 | 0 | 0.558 | 0 | 0.555 |
| 3069_IM-1432-10 | Pneumoni | 0 | 0.554 | 0 | 0.562 | 0 | 0.620 | 0 | 0.531 | 0 | 0.534 |
| 353_IM-1726-100 | Edema | 0 | 0.533 | 0 | 0.561 | 0 | 0.558 | 1 | 0.680 | 0 | 0.677 |
| 1262_IM-0178-00 | Pneumoth | 0 | 0.535 | 0 | 0.536 | 0 | 0.538 | 0 | 0.553 | 0 | 0.536 |

Per-image labels & scores (rows 16-29 of 29)

| filename | chexpert | Atelectasi | Atelectasi | Cardiome | Cardiome | Consolida | Consolida | Edema_gt | Edema_p | Pleural Ef | Pleural Effusion_p |
|-----------------|------------|------------|------------|----------|----------|-----------|-----------|----------|---------|------------|--------------------|
| 1262_IM-0178-10 | Pneumoth | 0 | 0.567 | 0 | 0.548 | 0 | 0.544 | 0 | 0.546 | 0 | 0.550 |
| 353_IM-1726-200 | Edema | 0 | 0.563 | 0 | 0.642 | 0 | 0.569 | 1 | 0.668 | 0 | 0.680 |
| 3348_IM-1605-20 | Cardiome | 0 | 0.558 | 1 | 0.749 | 0 | 0.551 | 0 | 0.627 | 0 | 0.530 |
| 541_IM-2147-100 | Fracture | 0 | 0.556 | 0 | 0.552 | 0 | 0.548 | 0 | 0.563 | 0 | 0.554 |
| 3348_IM-1605-10 | Cardiome | 0 | 0.542 | 1 | 0.755 | 0 | 0.544 | 0 | 0.581 | 0 | 0.527 |
| 3396_IM-1640-10 | Pleural Ef | 0 | 0.658 | 0 | 0.554 | 0 | 0.589 | 0 | 0.528 | 1 | 0.764 |
| 3069_IM-1432-20 | Pneumoni | 0 | 0.553 | 0 | 0.542 | 0 | 0.547 | 0 | 0.556 | 0 | 0.543 |
| 3638_IM-1804-20 | Cardiome | 0 | 0.548 | 1 | 0.641 | 0 | 0.551 | 0 | 0.573 | 0 | 0.544 |
| 3233_IM-1530-20 | Atelectasi | 1 | 0.592 | 0 | 0.561 | 0 | 0.559 | 0 | 0.546 | 0 | 0.554 |
| 1430_IM-0277-20 | Pneumoni | 0 | 0.581 | 0 | 0.560 | 0 | 0.550 | 0 | 0.576 | 0 | 0.619 |
| 2352_IM-0918-10 | Fracture | 0 | 0.545 | 0 | 0.545 | 0 | 0.541 | 0 | 0.534 | 0 | 0.656 |
| 2502_IM-1027-10 | Pleural Ef | 0 | 0.553 | 0 | 0.562 | 0 | 0.551 | 0 | 0.684 | 1 | 0.710 |
| 3233_IM-1530-10 | Atelectasi | 1 | 0.644 | 0 | 0.551 | 0 | 0.556 | 0 | 0.555 | 0 | 0.548 |
| 1262_IM-0178-20 | Pneumoth | 0 | 0.565 | 0 | 0.564 | 0 | 0.543 | 0 | 0.558 | 0 | 0.554 |